Hybridizing High Performance Computing, Distributed Systems, and Neural Networks for Optimal Investment Portfolio Management

Nahin Hossain
Computer Science and Engineering
Brac University
Dhaka, Bangladesh
nahin.hossain@g.bracu.ac.bd

Nawroze Baizid Mohammed
Computer Science and Engineering
Brac University
Dhaka, Bangladesh
nawroze.baizid.mohammed@g.bracu.ac.bd

Azmain Morshed
Computer Science
Brac University
Dhaka, Bangladesh
azmain.morshed@g.bracu.ac.bd

Mohammad Tajrian Arafat Alvi
Computer Science and Engineering
Brac University
Dhaka, Bangladesh
tajrian.arafat.alvi@g.bracu.ac.bd

Md Humaion Kabir Mehedi
Computer Science and Engineering
Brac University
Dhaka, Bangladesh
humaion.kabir.mehedi@g.bracu.ac.bd

Mehnaz Ara Fazal
Computer Science and Engineering
Brac University
Dhaka, Bangladesh
mehnaz.ara.fazal@g.bracu.ac.bd

Annajiat Alim Rasel
Computer Science and Engineering
Brac University
Dhaka, Bangladesh
annajiat@gmail.com

I. Related Work

A. HPC-Driven Weed Management in Precision Agriculture

Manuel López-Martíne [3] conducted an in-depth investigation into the application of high-performance computing (HPC) in precision agriculture (PA), focusing specifically on the challenging task of managing weed proliferation in crop fields. The paper effectively underscores the integration of sophisticated automatic methods and computational resources, particularly highlighting the role of HPC in the accurate identification, classification, and detection of diverse weed species. This innovative approach is essential for processing the extensive digital data collected from agricultural fields, significantly improving the precision and effectiveness of weed control strategies.

Central to this research is the strategic deployment of a high-performance computing cluster (HPCC), serving as a foundation for executing advanced image processing and analysis. This is accomplished through the application of cutting-edge deep learning techniques, notably convolutional neural networks (CNNs) utilizing the VGG16 and InceptionV3 models. The selection of these specific models is critical, as they demonstrate exceptional proficiency in distinguishing various weed species, thereby illustrating the substantial benefits of utilizing HPCC in PA for intricate image classification tasks. The study meticulously details the technical components of this approach, particularly the adoption of the Keras and Horovod frameworks within the HPCC to enable efficient distributed computing, crucial for the training of the CNN models. The research highlights the outstanding performance of the InceptionV3 model, which achieved an impressive processing time of 37 minutes and 55.193 seconds using six HPCC cores, along with a noteworthy accuracy rate of 0.65. This achievement is especially significant given the complexity of the computations and the size of the datasets involved.

Furthermore, the paper emphasizes the broader implications of these findings, suggesting that the integration of sophisticated computational techniques such as HPC and deep learning can greatly enhance weed management in precision agriculture. This advancement is instrumental in fostering more sustainable and effective agricultural practices, vital in the context of escalating global food demands and the necessity for eco-friendly farming techniques. Consequently, the outcomes of this study not only contribute significantly to the field of agricultural technology but also have extensive ramifications for the future of sustainable agriculture practices.

B. Artificial Intelligence in Asset Management

Bartram's paper titled Artificial Intelligence in Asset Management presents an extensive exploration of the integration of artificial intelligence (AI) within the financial sector, with a particular focus on the field of asset management. This in-depth analysis encompasses the profound impact of AI across critical domains, including portfolio management, trading strategies, and risk management.

Throughout the paper, the authors delve into the sophisticated AI methodologies at play, spotlighting the significance of ML, NN, and NLP. These advanced AI techniques are showcased for their remarkable ability to revolutionize traditional practices within asset management, yielding substantial improvements in market risk prediction, portfolio optimization, and the development of astute trading strategies.

In addition to highlighting the transformative potential of AI, the paper offers a pragmatic examination of the challenges inherent in its integration. The complexities surrounding model development and the imperative reliance on comprehensive and high-quality data sources are thoughtfully discussed.

Ultimately, this research underscores the dynamic evolution taking place within the financial industry, as it steadily gravitates towards automated, data-driven methodologies. It recognizes the immense promise that AI holds for reshaping asset management practices, while simultaneously acknowledging the intricacies and potential risks associated with these cutting-edge technologies. This paper serves as a pivotal resource for those seeking to comprehend the profound implications of AI adoption within the financial landscape, providing invaluable insights into both the opportunities and complexities that lie ahead.

C. Machine Learning

Pawel Gepner [2] investigated the use of advanced machine learning algorithms in optimizing HPC system performance. Their approach focused on efficient resource allocation and power management to achieve higher computational throughput without significant increases in power consumption. Another study by a researcher explored the integration of AI techniques in HPC systems, demonstrating how machine learning models can enhance the accuracy and efficiency of complex simulations typically run on HPC platforms. Additionally, the author provided insights into the architectural innovations in hardware design tailored to support the intensive computational requirements of both HPC and AI applications, suggesting a synergistic approach to system development.

These studies collectively underscore the importance of a multidisciplinary approach in tackling the challenges posed by the next generation of HPC systems. The fusion of machine learning, AI techniques, and innovative hardware designs appears to be a promising pathway towards achieving Exascale computing while managing power and cost constraints.

D. Spiking Neural Networks

Rong Zhao [1] propose a novel framework for the design and computation of Hybrid Neural Networks (HNNs) by integrating Spiking Neural Networks (SNNs) and Artificial Neural Networks (ANNs). This integration is facilitated through the introduction of Hybrid Units, which serve as a connection interface between these two computing paradigms. The proposed framework not only amalgamates the key features of SNNs and ANNs but also enhances flexibility and efficiency by decoupling these features. The HUs efficiently enable the transmission and modification of information within HNNs because of their designable and learnable nature. The efficacy of this framework is demonstrated through three distinct case studies. The first, a hybrid sensing network, achieves high tracking accuracy and energy efficiency by implementing multi-pathway sensing. The second, a hybrid modulation network, enables meta-continual learning across multiple tasks through hierarchical information abstraction. Finally, the hybrid reasoning network showcases robust, interpretable, and parallel multimodal reasoning. This research by the author marks a significant stride in cross-paradigm modeling, broadening the scope for intelligent task applications in various fields.

II. Methodology

A. Data Acquisition with yfinance

The research utilized yfinance, a Python library that provides a seamless method to access Yahoo Finance's historical financial data. It was selected for its ability to programmatically download extensive stock market information, including stock prices, dividends, and financial summaries. This automated access is vital for performance analysis, backtesting trading strategies, and developing predictive models. In our study, yfinance facilitated the acquisition of Apple Inc.'s historical stock prices for our analysis. We reshaped the closing prices from -1 to 1 and data preprocessing was not done as we used the data from yfinance library.

B. LSTM Architecture and Application

Long Short-Term Memory networks are a key element in deep learning for handling sequential data, such as time series analysis crucial in financial forecasting. The architecture of LSTM is designed to overcome the vanishing gradient problem seen in traditional Recurrent Neural Networks (RNNs), enabling it to effectively capture long-term dependencies. The input, output, and forget gates are arranged in a complicated way across the layers of the LSTM network in our study. The forget gate processes the current input and the previous concealed state using a sigmoid function to determine what data should be discarded from the cell state. The candidate memory is a filtered representation of the cell state, and the input gate controls how the cell state is updated. The output of the forget gate is combined with the result of the input gate and candidate memory to update the previous cell state to the current one. The following hidden state, which contains information about earlier inputs, is determined by the output gate.

In our research, we built a multi-layer LSTM model to enhance learning. The first LSTM layer had 50 units, returning the full sequence to the next layer, and the input shape was set according to the input data's dimensions. A second LSTM layer was added, outputting a 2D shape suitable for a Dense layer. Dense layers transitioned the LSTM's output to the final prediction. The LSTM model was trained using the backpropagation through time (BPTT) algorithm on standardized stock price data, scaled between 0 and 1. This standardization is vital for the effective operation of gradient descent optimization. The model's weights were refined through iterative training epochs, minimizing the Mean Squared Error (MSE) to improve predictive accuracy. The training process balanced model performance and efficiency by optimizing batch size and the number of epochs.

C. Linear Regression Implementation

Linear Regression, a foundational statistical method, aims to model the relationship between a dependent variable and one or more independent variables by fitting a linear equation to observed data. The simplicity of the model lies in its assumption of a linear association between variables, making it a versatile choice for straightforward relationships. However, when confronted with the intricate and non-linear dynamics inherent in stock market data, Linear Regression may encounter limitations in capturing complex patterns.

D. Train Test Splitting

The data is divided into an 80% training and 20% testing split. Each data point has 60 input features related to financial indicators, and the 61st feature is the target label for investment outcomes. The neural network trains on the 80% portion to understand market trends and tests its predictions on the remaining 20% to ensure accuracy and reliability in real-world scenarios. This process, repeated across the dataset, continually enhances the system's effectiveness in portfolio management.

III. Experimental Results and Analysis

A. LSTM Performance Evaluation

The Long Short-Term Memory (LSTM) model exhibited a robust ability to forecast Apple Inc.'s stock closing prices. By learning from sequences of historical price data, the LSTM could capture complex patterns essential for accurate predictions. The visualization of the model's performance depicted the predicted versus actual closing prices, showcasing a high degree of alignment, especially towards the end of the timeline. The overlapping lines in the graph indicate that the LSTM predictions closely mirrored the actual prices, capturing both the general upward trend and the fluctuations within the market. This is quantitatively supported by a Root Mean Squared Error (RMSE) of approximately 77.33 units, reflecting the average deviation of the predictions from the actual values.

B. Linear Regression Analysis

In stark contrast, Linear Regression struggled to model the same stock price data effectively. The simplicity of Linear Regression, which assumes a linear relationship between input and output variables, proved inadequate for the complexity inherent in financial time series data. This was demonstrated by a high Mean Squared Error (MSE) and a notably large RMSE of greater than 101, which indicated that the model's predictions were frequently inaccurate. Furthermore, because the model was unable to detect any underlying patterns or trends, a negative R-squared value suggested that it was inappropriate for the dataset.

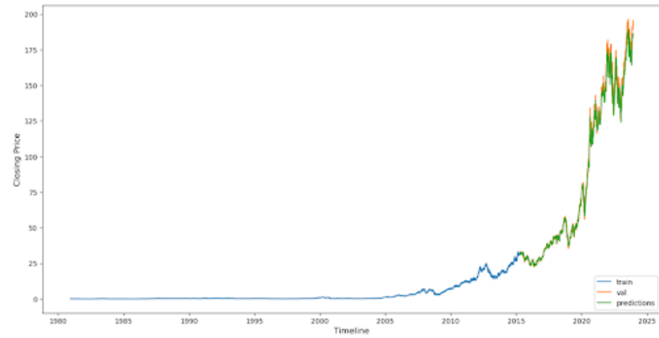

Fig. 1: Long-Term Trend of Stock Prices from 1980 to 2025

Fig. 2: Short-Term Fluctuations in Stock Prices from 2015 to 2024

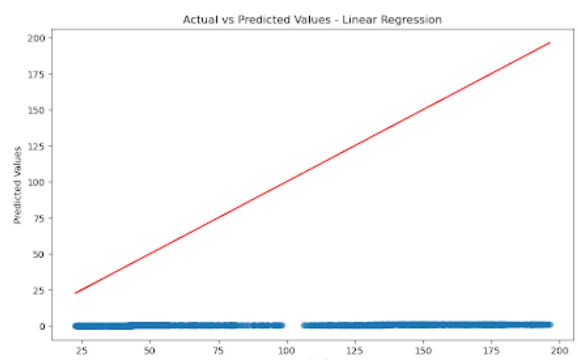

Fig. 3: Comparison of Actual Values to Predictions by Linear Regression Model

C. Model Comparison

The comparison between the Long Short-Term Memory (LSTM) model and Linear Regression in forecasting Apple

Inc.'s stock closing prices reveals stark differences in performance. The LSTM model, leveraging its ability to capture intricate patterns from historical price data sequences, demonstrated robustness by closely aligning its predictions with actual prices. Its visualization depicted a high level of agreement, capturing both the overarching upward trend and the market fluctuations. Quantitatively, the LSTM exhibited an average Root Mean Squared Error (RMSE) of approximately 77.33 units, indicating minimal deviation from actual values. In contrast, Linear Regression struggled significantly with this task due to its oversimplified assumption of a linear relationship between input and output variables. The model's high Mean Squared Error (MSE) and RMSE exceeding 101 underscored its frequent inaccuracies and inability to capture the complexity of financial time series data. Additionally, the negative R-squared value further emphasized its unsuitability for analyzing this dataset. Overall, the LSTM's adaptability to intricate patterns and nuanced trends outperformed the limitations of Linear Regression in forecasting stock prices accurately.

D. GPU versus CPU Training Times

The use of TensorFlow GPU in Google Colab for training the LSTM model resulted in a significant reduction in training time compared to a CPU-based setup. GPU acceleration proved to be a game-changer, enabling faster computation and more efficient training of the complex LSTM network. The speed and performance gains from using a GPU were evident, highlighting the importance of computational resources in conducting high-performance computing for deep learning tasks.

E. Comparative Analysis of Epochs

The LSTM model's training epochs, visualized on two different backgrounds, illustrate the stark difference in performance between GPU and CPU training sessions. The GPU sessions, shown on a black background, completed faster and displayed a steady decrease in loss over epochs, signifying a more efficient and effective training process. In contrast, the CPU sessions, represented on a white background, were slower and less consistent, emphasizing the GPU's superior computational capabilities

IV. Future Plans and Challenges

Future research in stock market forecasting will focus on employing advanced machine learning techniques like Random Forest, Gradient Boosting, RNNs, and LSTMs, which are better suited for complex, non-linear data. This approach will be complemented by enriching the dataset with a wider array of variables, including market sentiment and economic indicators, and applying robust cross-validation and time-series analysis methods for improved accuracy and reduced overfitting. The primary challenge lies in balancing the complexity of advanced models with their interpretability, ensuring they are understandable and not just accurate. Another significant hurdle is maintaining high-quality data in the ever-changing financial market. Computational demands, especially for real-time processing in resource-intensive models, also pose a major challenge in developing effective and insightful stock market forecasting models.

V. Conclusion

This research introduces an innovative framework that synergizes the computational power of High Performance Computing (HPC), the scalability of distributed systems, and the predictive capabilities of neural networks to advance investment portfolio management. Our approach addresses the complexities and dynamic nature of the global financial markets with a system designed to enhance efficiency and market responsiveness. Through the application of sophisticated neural network architectures and distributed computing strategies to analyze extensive financial datasets, we have observed marked improvements in portfolio performance metrics over traditional methods. The proposed method is rigorously evaluated for accuracy, responsiveness, and scalability, proving to be a formidable new standard in the realm of financial portfolio optimization and management. The adoption of this method offers a robust and scalable solution, with performance and efficiency that make it a highly viable tool for real-time financial analysis and decision-making.

REFERENCES

[1] Zhao, R., Yang, Z., Zheng, H. et al. A framework for the general design and computation of hybrid neural networks. Nat Commun 13, 3427 (2022). https://doi.org/10.1038/s41467-022-30964-7

[2] Gepner, Pawel. (2021). Machine Learning and High-Performance Computing Hybrid Systems, a New Way of Performance Acceleration in Engineering and Scientific Applications. 27-36. 10.15439/2021F004.

[3] López-Martínez, M.; Díaz-Flórez, G.; Villagrana-Barraza, S.; Solís-Sánchez, L.O.; Guerrero-Osuna, H.A.; Soto-Zarazúa, G.M.; Olvera-Olvera, C.A. A High-Performance Computing Cluster for Distributed Deep Learning: A Practical Case of Weed Classification Using Convolutional Neural Network Models. Appl. Sci. 2023, 13, 6007. https://doi.org/10.3390/app13106007.

[4] Abe, Masaya, and Hideki Nakayama. 2018. "Deep Learning for Forecasting Stock Returns in the Cross-Section." In Advances in Knowledge Discovery and Data Mining, Part 1, edited by Dinh Phung, Vincent S. Tseng, Geoffrey I. Webb, Bao Ho, Mohadeseh Ganji, and Lida Rashidi, 273–84. Cham, Switzerland: Springer International Publishing.

[5] Aggarwal, Charu C. 2018. Neural Networks and Deep Learning. Cham, Switzerland: Springer. https://doi.org/10.1007/978-3-319-94463-0.

[6] Aggarwal, Charu C., and Chandan K. Reddy, eds. 2014. Data Clustering: Algorithms and Applications. Boca Raton, FL: CRC Press

[7] Ahmed, Nesreen K., Amir F. Atiya, Neamat El Gayar, and Hisham El-Shishiny. 2010. "An Empirical Comparison of Machine Learning Models for Time Series Forecasting." Econometric Reviews 29 (5–6): 594–621. https:// doi.org/10.1080/07474938.2010.481556.

[8] Ahn, Jae Joon, Kyong Joo Oh, Tae Yoon Kim, and Dong Ha Kim. 2011. "Usefulness of Support Vector Machine to Develop an Early Warning System for Financial Crisis." Expert Systems with Applications 38 (4): 2966–73. https:// doi.org/10.1016/j.eswa.2010.08.085.

[9] Alberg, John, and Zachary C. Lipton. 2017. "Improving Factor-Based Quantitative Investing by Forecasting Company Fundamentals," version 2. arXiv.org/abs/1711.04837v2.